

Introduction To Deep Learning & Neural Networks with keras (M2) (2)

Created	@December 10, 2025 7:12 PM
Module Code	IBMM02: Introduction to Deep Learning

Module 4 : Deep Learning Models

Shallow vs Deep Neural Networks

Type	Layers	What it handles	Use cases
Shallow Neural Network	1-2 hidden layers	Input must be preprocessed into numeric vectors	Simple problems, structured data
Deep Neural Network	3+ hidden layers (often many more)	Can take raw data (images, text, audio) and learn features automatically	Computer vision, NLP, speech, advanced ML tasks

Convolutional Neural Networks

A **CNN** is a special type of neural network designed specifically for **images**.

A normal neural network takes this kind of input:

```
[ x1, x2, x3, x4, ... ]
```

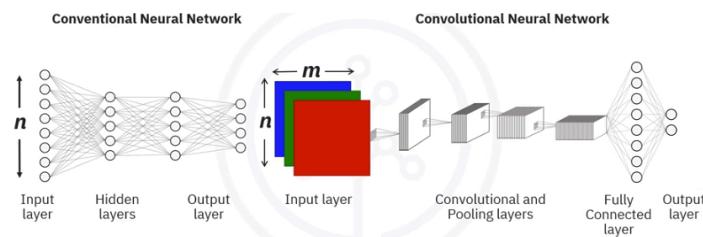
But images are **2D (or 3D)**:

Height × Width × Channels (RGB)

So CNNs **keep the spatial structure** of images instead of flattening everything at the start.

If we flatten an image (e.g., $128 \times 128 \times 3 \rightarrow 49,152$ numbers) and feed it into a fully connected network:

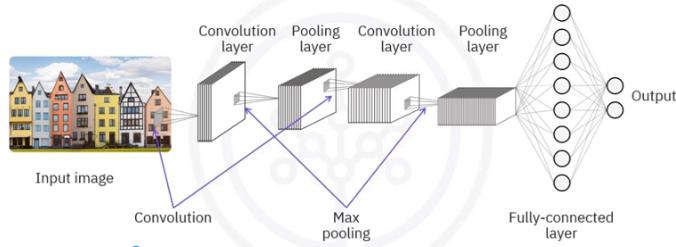
- Too many parameters → **slow training, high memory usage**
- Very easy to **overfit**



CNN Architecture

A CNN usually follows this order:

- 1 Convolution Layer**
- 2 ReLU Activation**
- 3 Pooling Layer**
(repeat 1–3 several times)
- 4 Flatten**
- 5 Fully Connected (Dense) Layers**
- 6 Output Layer (Softmax)**



1. Convolution Layer — Feature extractor

This layer applies **filters (kernels)** to the image.

A filter is a tiny matrix, like:

$$\begin{bmatrix} 1 & 0 \\ -1 & 1 \end{bmatrix}$$

The filter slides over the image → performs dot product → creates a **feature map**.

This helps the network learn:

- Edges
- Corners
- Colors
- Patterns
- Shapes

Using convolution dramatically **reduces number of parameters**.

Each filter = new feature map.

Using more filters → detect more types of features.

2. ReLU Activation

ReLU simply means:

```
if x < 0 → 0
else → x
```

This introduces **non-linearity** and makes training faster/less likely to have vanishing gradients.

3. Pooling Layer — Reduces Size

Pooling compresses image size:

Max Pooling (most common)

Keeps only the maximum value in each region.

```
[1, 3]  
[2, 5] → max = 5
```

Why?

- Simpler computation
- Reduces overfitting
- Makes model recognize objects even if position changes a bit → **spatial invariance**

Average Pooling

Takes the average of values in each region.

4. Fully Connected Layer

Once features are extracted, they are flattened into a long vector (1D) :

```
[ f1, f2, f3, ... ]
```

Then fed into:

- Dense layer(s)

- Output layer (softmax for classification)

This part works just like a normal neural network.

When are CNNs used?

- ✓ Image classification
- ✓ Object detection
- ✓ Medical imaging
- ✓ Face recognition
- ✓ Traffic sign detection
- ✓ OCR (text from images)
- ✓ Satellite image analysis

CNNs handle **raw images** without manual feature engineering.

```
model = Sequential()
model.add(Conv2D(16, (2,2), activation='relu', input_shape=(128,128,3)))
model.add(MaxPooling2D(pool_size=2))

model.add(Conv2D(32, (2,2), activation='relu'))
model.add(MaxPooling2D(pool_size=2))

model.add(Flatten())
model.add(Dense(100, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))

"""16 filters → then 32 filters
2×2 pooling
Dense layer with 100 units
Final softmax output"""
```

Recurrent Neural Networks

Traditional neural networks assume:

- | Each input is independent.

But many real-world problems involve **sequences**, where the order matters, for example:

- Words in a sentence
- Scenes in a movie
- Stock prices over time
- DNA sequences
- Handwriting strokes

A normal neural network cannot understand sequence order.

RNNs solve this.

1. What Problem Do RNNs Solve?

RNNs are built to handle **sequential data**.

They add a **memory** mechanism, meaning:

- ✓ They take the current input
- ✓ AND they also use information from previous inputs

This lets the network understand **context** and **temporal patterns**.

2. How RNNs Work (Architecture)

Imagine a normal neural network, but with a loop:

At time **t = 0**

- Input: x_0
- Output: a_0

At time **t = 1**

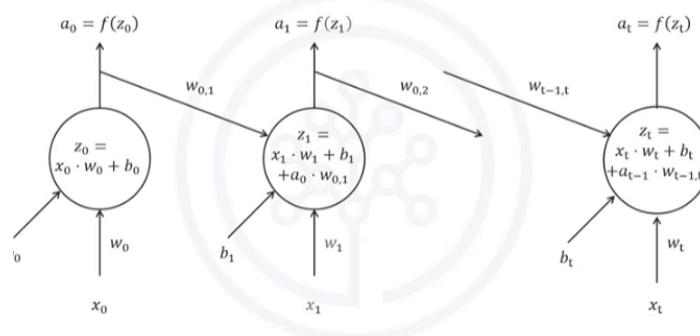
- Because it's sequential, the network takes:

- New input: x_1
- Previous output: a_0

So at each step:

$$\text{output}_t = f(\text{input}_t + \text{previous_output}_{t-1})$$

This “loop” lets information flow from previous time steps.



3. What Are RNNs Used For?

RNNs are good at anything involving **time, order, or dependencies**.

Common uses:

- Text processing
- Language modeling
- Predicting next word
- Genomics
- Handwriting recognition
- Speech recognition
- Stock market prediction
- Audio generation

Basically anything where the meaning depends on what came before.

4. The LSTM (Long Short-Term Memory)

Basic RNNs have a problem:

They **forget long-term information** and suffer from **vanishing gradients**.

LSTMs fix this.

An LSTM has special **gates** that control:

- What to remember
- What to forget
- What to output

So LSTMs can capture **long-range dependencies** (like remembering the subject of a sentence many words later).

5. Real-World Applications of LSTMs

LSTMs are widely used for:

✓ Image generation

Using sequential pixel patterns

✓ Handwriting generation

Reproducing realistic handwritten strokes

✓ Automatic image captioning

CNN extracts features → LSTM generates sentence

✓ Video description

LSTM learns sequences of visual features over time

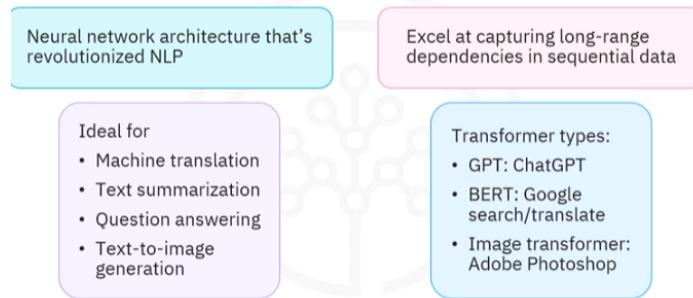
They are extremely powerful in tasks involving long sequences.

Transformers

Transformers are a type of deep learning architecture that changed the entire field of **NLP, vision, and generative AI**. They are used in today's powerful AI systems—like ChatGPT, Gemini, BERT, and image generators such as DALL·E.

The key reason they became dominant:

Transformers can capture long-range relationships in data extremely well and can be trained in parallel, making them fast and scalable.



1. Why Transformers?

Before transformers:

- **RNNs** could handle sequences but were slow (trained step by step) and struggled with long dependencies.
- **CNNs** could understand patterns but were not built for text or long sequences.

Transformers solve both problems:

They understand **context across the entire sequence**, not just nearby words

They process all tokens **in parallel**, making training **much faster**

2. The Core Idea: Attention

Transformers rely on a mechanism called **attention**.

Instead of reading tokens one by one, the model:

- Looks at the whole sequence at once
- Decides which words are important for understanding each token

Example:

In the sentence "*The dog ran because it was scared,*"

the word "**it**" should be linked to "**dog**".

Attention makes this possible.

3. Self-Attention: How It Works

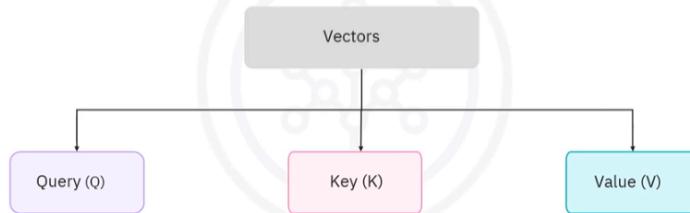
Self-attention is used when processing text.

For each token, the model generates:

- **Query (Q)** → the token we are focusing on
- **Key (K)** → signals how relevant each token is
- **Value (V)** → the information carried by each token

Self-attention mechanism

- Transformers use self-attention mechanism to process text data
- Self-attention mechanism has three parts:



Steps:

1. **Q, K, V vectors are created** for every token
2. The model computes **attention scores** → $Q \cdot K$
3. Scores are scaled and passed through **Softmax** to get probabilities
4. The model computes a **weighted sum of V values**
→ this produces a *contextualized meaning* of the token

This process helps each word understand the entire sentence.

Self-attention mechanism example

Input	The	Dog	Runs
Embeddings 1 0 1 0	0 2 0 2	1 1 1 1	
Positional Encoding 2 1 2 1	2 4 2 4	4 4 4 4	

Sequence of words:
"The dog runs"

Each word in sequence is represented by a vector

Positional information is added to each word

Self-attention mechanism example (continue)

We generate queries (Q), keys (K), and values (V)

Input	The	Dog	Runs
Embeddings	1 0 1 0	0 2 0 2	1 1 1 1
Positional Encoding	2 1 2 1	2 4 2 4	4 4 4 4
Queries	Q1 1 0 2	Q2 2 2 2	Q3 2 1 3
Keys	K1 0 1 1	K2 4 4 0	K3 2 3 1
Values	V1 1 2 3	V2 2 8 0	V3 2 6 3

4. Cross-Attention: For Text-to-Image Generation

Cross-attention is used when **one type of data guides another**, such as:

- Text → influences image generation

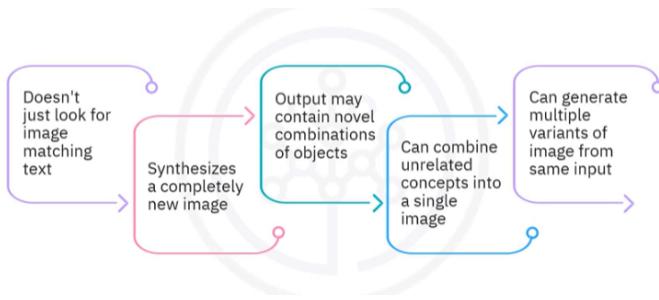
Example:

Prompt: "a two-story house with a red roof"

Process:

1. The text is converted into embeddings
2. The transformer's encoder produces a sequence of Q (from text)
3. The image-generation model (e.g., DALL-E) uses cross-attention so image tokens respond to text tokens
4. The model autoregressively generates image patches based on:
 - Text meaning
 - Previously generated image parts

This allows the model to **synthesize new images**, not just retrieve existing ones.



5. Power of Transformers

Transformers are used in:



- Translation
- Summarization
- Answering questions
- Chatbots



- Text-to-image generation
- Style transfer
- Novel concept generation



- Search (Google uses BERT)
- Photoshop's image editing features
- Code generation (Copilot)

They handle extremely long and complex patterns better than older models.

6. Limitations of Transformers

Despite their power, transformers have some drawbacks:

- ! They require huge datasets

Without massive training data, they don't generalize well.

! They inherit bias from training data

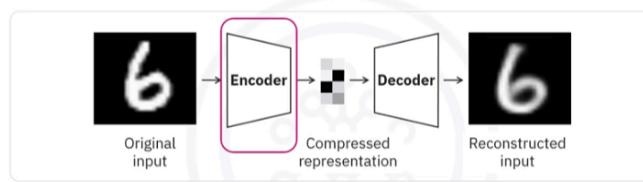
If the data contains stereotypes or errors, the model can learn them.

! They need large compute resources

Transformers are computationally expensive.

Autoencoder

Autoencoder architecture



What are Autoencoders?

Autoencoders are **unsupervised neural networks** used mainly for:

- **Compressing data** (encoding)
- **Reconstructing data** (decoding)

They learn to recreate the input **as output**, so the target = input.

Think of them as a machine that learns:

"How can I represent this data in a smaller form... and bring it back?"

How They Work (Architecture)

1. Encoder

- Takes the input (image/text/etc.)
- Compresses it into a smaller representation (called *latent vector*)

2. Bottleneck

- The compressed "code"

- Forces the network to keep only the most important features

3. Decoder

- Expands the compressed representation back to original shape

Training Objective

Minimize:

Original Input \approx Reconstructed Output

So, the network learns meaningful patterns.

Why Autoencoders are Useful

1. Data Denoising

Feed a noisy image \rightarrow autoencoder learns to remove noise.

2. Dimensionality Reduction

Similar to PCA, but **non-linear**, so more powerful.

3. Feature Learning

Automatically extracts patterns in data.

Restricted Boltzmann Machines (RBMs)

RBM s are a special type of autoencoder-like model (older, but powerful):

Applications:

✓ Fixing **imbalanced datasets**

- RBMs learn the distribution of the minority class
- Then they **generate synthetic samples**
- Helps balance class counts

✓ Estimating missing values

- RBMs learn correlations in features
- Can fill missing entries

✓ Automatic feature extraction

- Especially useful for unstructured data (e.g., images)
-

Using Pretrained Model

What is a Pretrained Model?

A **pretrained model** is a neural network that has already been trained on a large dataset (like **ImageNet**, with 14M images / 1000 classes).

Examples:

- **VGG16**
- **ResNet**
- **MobileNet**
- **EfficientNet**

These models have already learned:

- Edges
- Textures
- Shapes
- Object parts

...so you don't need to train them from scratch.

Two Ways to Use Pretrained Models

1. Use as a Fixed Feature Extractor

- You **freeze** all layers → no retraining.
- Pass new images through the network.

- The model outputs **feature maps** (high-level patterns).
- You use these features for:
 - Clustering
 - Visualization
 - Dimensionality reduction
 - Feeding into ML models (SVM, RF, etc.)

Benefits

- ✓ No training needed
- ✓ Very fast
- ✓ Works well with small datasets
- ✓ Uses very powerful learned features

Limitations

- ✗ Cannot adapt to your dataset
- ✗ Performance limited if your data is very different from ImageNet

```

import os
import shutil
from PIL import Image
import numpy as np

# Define the base directory for sample data
base_dir = 'sample_data'
class1_dir = os.path.join(base_dir, 'class1')
class2_dir = os.path.join(base_dir, 'class2')

# Create directories for two classes
os.makedirs(class1_dir, exist_ok=True)
os.makedirs(class2_dir, exist_ok=True)

# Function to generate and save random images
def generate_random_images(save_dir, num_images):
    for i in range(num_images):
        # Generate a random RGB image of size 224x224

```

How to use pretrained models as feature extractors in Keras

```

# Generate a random RGB image of size 224x224
img = Image.fromarray(np.uint8(np.random.rand(224, 224, 3) * 255))
# Save the image to the specified directory
img.save(os.path.join(save_dir, f'image_{i}.jpg'))

# Number of images to generate for each class
num_images_per_class = 100 # You can increase this to have more training data

# Generate random images for class 1 and class 2
generate_random_images(class1_dir, num_images_per_class)
generate_random_images(class2_dir, num_images_per_class)

print(f'Sample data generated at {base_dir} with {num_images_per_class} images per class.')

```

Example: Pretrained model for extracting features

```
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam

# Load the VGG16 model pre-trained on ImageNet
base_model = VGG16(weights='imagenet'
, include_top=False, input_shape=(224, 224, 3))

# Freeze all layers initially
for layer in base_model.layers:
    layer.trainable = False

# Create a new model and add the base model and new layers
model = Sequential([
    base_model,
    Flatten(),
```

Example: Pretrained model for extracting features

```
# Create a new model and add the base model and new layers
model = Sequential([
    base_model,
    Flatten(),
    Dense(256, activation='relu'),
    Dense(1, activation='sigmoid')

# Change to the number of classes you have
])

# Compile the model
model.compile(optimizer=Adam(learning_rate=0.001),
loss='binary_crossentropy', metrics=['accuracy'])
```

Example: Pretrained model for extracting features

```
# Load and preprocess the dataset
train_datagen = ImageDataGenerator(rescale=1./255)
train_generator = train_datagen.flow_from_directory(
    '/content/sample_data',
    target_size=(224, 224),
    batch_size=32,
    class_mode='binary'
)

# Train the model with frozen layers
model.fit(train_generator, epochs=10)

# Gradually unfreeze layers and fine-tune
for layer in base_model.layers[-4:]: # Unfreeze the last 4 layers
    layer.trainable = True

# Compile the model again with a lower learning rate for fine-tuning
model.compile(optimizer=Adam(learning_rate=0.0001),
loss='binary_crossentropy', metrics=['accuracy'])

# Fine-tune the model
model.fit(train_generator, epochs=10)
```

2.Fine-Tuning (Transfer Learning)

- You **unfreeze some top layers** of the pretrained model.
- Add new layers for your task.
- Train everything together (usually with low learning rate).

Why Fine-Tune?

Because the pretrained model may not perfectly match your new dataset.

For example:

- ImageNet → dogs, cars, buildings
- Your dataset → retinal fundus images

So fine-tuning helps the model adjust.

Benefits

- ✓ Much better accuracy
- ✓ Model adapts to your domain
- ✓ Still requires *far less* data than training from scratch

Keras Workflow (Conceptual Steps)

Step 1: Load pretrained model

```
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(2  
24, 224, 3))
```

Step 2: Freeze its layers

```
base_model.trainable = False
```

Step 3: Add your own classifier

Flatten → Dense(256, relu) → Dense(1, sigmoid)

Step 4: Train your model (feature extraction mode)

This only trains the new layers you added.

Step 5: (Optional) Fine-tune

Unfreeze top layers from VGG16 and train again with smaller LR.

When Should You Use Feature Extraction vs Fine-Tuning?

Scenario	Best Choice
Very small dataset	Feature extractor
Low compute	Feature extractor
Dataset very different from ImageNet	Fine-tuning
Want highest accuracy	Fine-tuning
Quick prototype	Feature extractor