



# Machine Learning with Python(M1)

|             |  |
|-------------|--|
| Created     | @December 9, 2025 10:50 AM               |
| Module Code | IBMM01: Introduction to Machine Learning |

## Module 5: Evaluating and Validating Machine Learning Models

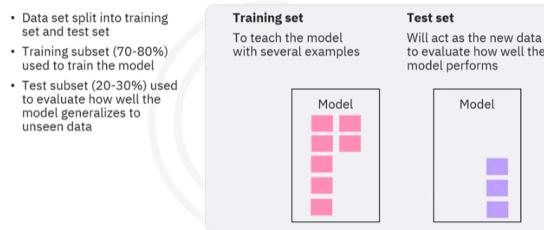
### 1. Supervised Learning Evaluation

- Purpose: Measures how well a machine learning model predicts outcomes on **unseen data**.
- Importance: Helps understand model effectiveness during both training and testing phases.
- Method: Compare **model predictions** to **ground truth labels**.

### 2. Train-Test-Split Technique

- Splits the dataset into:
  - Training set** (70–80%): Used to train the model.
  - Test set** (20–30%): Used to evaluate generalization to new data.
- Ensures the model is not overfitting to all available data.

#### The train/test split technique



### 3. Key Classification Metrics

#### 1. Accuracy

- Formula:

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

- Example: 70% accuracy if 7 out of 10 predictions are correct.

## 2. Confusion Matrix

- A table comparing **true labels** vs **predicted labels**.
- Components:
  - **True Positive (TP)**: Predicted positive, actually positive.
  - **True Negative (TN)**: Predicted negative, actually negative.
  - **False Positive (FP)**: Predicted positive, actually negative.
  - **False Negative (FN)**: Predicted negative, actually positive.

## 3. Precision

- Measures the accuracy of positive predictions.
- Formula:

$$\text{Precision} = \frac{TP}{TP + FP}$$

- Important when **false positives are costly**, e.g., movie recommendations.

## 4. Recall

- Measures coverage of actual positives.
- Formula:

$$\text{Recall} = \frac{TP}{TP + FN}$$

- Important when **false negatives are costly**, e.g., medical diagnoses.

## 5. F1 Score

- Harmonic mean of precision and recall:

$$F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

- Useful when **both precision and recall are important**, e.g., healthcare scenarios.

## 4. Practical Example

- **Pass/Fail Test Predictions:**
  - Correct predictions = green squares.
  - Incorrect predictions = grey squares.
  - Accuracy = Correct / Total.
  - Precision & recall calculated specifically for the "pass" class.
- **Iris Flower Classification with KNN:**
  - Accuracy: 93%.
  - Confusion matrix visualized with a heatmap.
  - Setosa class achieved perfect precision, recall, and F1 score.

## 5. Summary Table

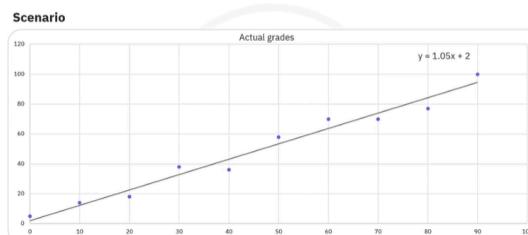
| Metric    | Purpose                              | Formula  | Use Case Example                   |
|-----------|--------------------------------------|--|------------------------------------|
| Accuracy  | Overall correctness                  | $(TP + TN) / \text{Total}$                                       | General model performance          |
| Precision | Correctness of positive predictions  | $TP / (TP + FP)$   | Movie recommendations              |
| Recall    | Coverage of actual positives         | $TP / (TP + FN)$   | Medical diagnosis                  |
| F1 Score  | Balance between precision and recall | $2 * (\text{Precision} * \text{Recall}) / (\text{P} + \text{R})$ | Health predictions, critical tasks |

## Regression Metrics and Evaluation Techniques

### 1. Purpose of Regression Evaluation

- Regression models predict **continuous numerical values** (e.g., exam scores, house prices).
- Evaluation is needed because models **can make errors**, and measuring these errors shows how well a model predicts unseen data.

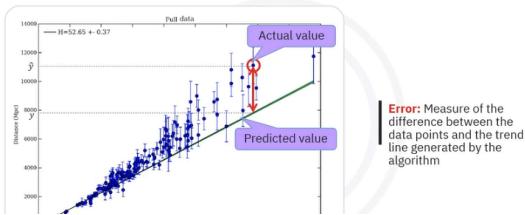
#### Evaluating regression models



Error = Difference between the predicted value ( $\hat{y}_i$ ) and the actual value ( $y_i$ ):

$$\text{Error}_i = \hat{y}_i - y_i$$

## Error of the model



## 2. Key Regression Metrics

### 1. Mean Absolute Error (MAE)

- Average of absolute differences between predicted and actual values.
- Formula:

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |\hat{y}_i - y_i|$$

- Advantages: Easy to interpret, same units as the target variable.

### 2. Mean Squared Error (MSE)

- Average of **squared differences** between predicted and actual values.
- Formula:

$$\text{MSE} = \frac{1}{n-p} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

- Penalizes larger errors more heavily.

### 3. Root Mean Squared Error (RMSE)

- Square root of MSE.
- Formula:

$$\text{RMSE} = \sqrt{\text{MSE}}$$

- Popular because it has the **same units** as the target variable, making it easy to interpret.

### 4. R-squared (R<sup>2</sup>)

- Proportion of variance in the target variable explained by the model.

- Formula:

$$R^2 = \frac{\sum_{i=1}^n (\hat{y}_i - \bar{y})^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

$$R^2 = 1 - \frac{\text{Unexplained Variance}}{\text{Total Variance}}$$

- Range:
    - 1 → Perfect fit
    - 0 → Model explains nothing beyond the mean
    - <0 → Model performs worse than predicting the mean
  - Note: Assumes a **linear relationship**; can be misleading for nonlinear models.
- 

### 3. Explained Variance

#### 3. Explained Variance

- Measures how much of the **total variability** in the target is captured by the model.
- Perfect predictor: Explained variance = Total variance, Unexplained variance = 0 →  $R^2 = 1$
- Simplistic predictor (predicts mean for all points): Explained variance = 0 →  $R^2 = 0$

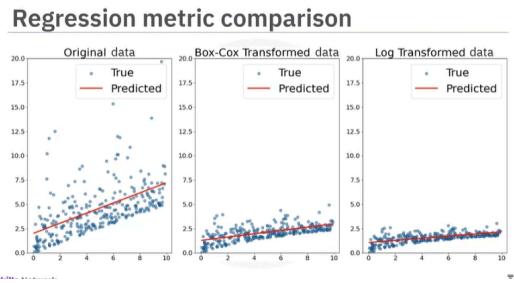
#### Explained variance and $R^2$

$$\begin{aligned}\text{Explained Variance} &= \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 \\ \text{Unexplained Variance} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 = n * \text{MSE} \\ \text{Total Variance} &= \sum_{i=1}^n (y_i - \bar{y})^2 = \sigma^2 \\ R^2 &= \frac{\sum_{i=1}^n (\hat{y}_i - \bar{y})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} = 1 - \frac{\text{Unexplained Variance}}{\text{Total Variance}}\end{aligned}$$

### 4. Model Transformation Examples

- For skewed data (e.g., log-normal distribution):
  - Transformations like **Box-Cox** or **log transformation** can improve linear model fit.
  - Metrics after transformation show:
    - Higher  $R^2$
    - Lower MAE, MSE, RMSE

- Visual inspection confirms better alignment with the trend line.



## 5. Key Takeaways

- **Error metrics (MAE, MSE, RMSE)** quantify the magnitude of prediction errors.
- **R-squared** indicates how well variance in the target is explained by the model.
- **Visualizing predictions vs actuals** is important for understanding model fit.
- Transformations can significantly improve regression performance, especially with skewed data.

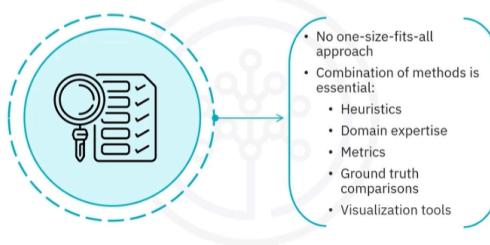
# Evaluating Unsupervised Learning Models: Heuristics and Techniques

## 1. Challenges in Evaluating Unsupervised Learning

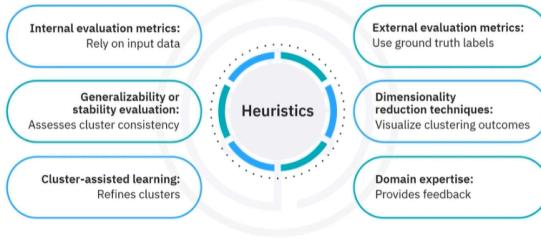
- **No predefined labels:** Unlike supervised learning, there's no ground truth to directly compare predictions.
- **Subjectivity:** Evaluating patterns or clusters often requires domain expertise and careful interpretation.
- **Stability:** Ensures consistent results across different subsets or perturbations of the data.

## 2. Evaluation Approaches

### Why is evaluation critical?



## Important clustering heuristics



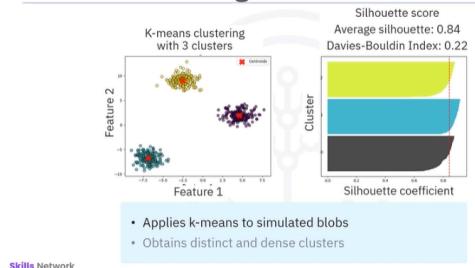
### 1. Internal Metrics – Use only the input data to assess clustering quality:

- **Silhouette Score:**

Measures cohesion vs separation, ranges -1 to 1; higher = better-defined clusters.

- **Davies-Bouldin Index:** Ratio of intra-cluster compactness to inter-cluster separation; lower = better clusters.
- **Inertia (K-means):** Sum of squared distances within clusters; lower = more compact clusters, but tradeoff exists with more clusters.

## Internal clustering evaluation



### 2. External Metrics – Require ground truth labels (if available):

- **Adjusted Rand Index (ARI):** Compares clustering to true labels, -1 to 1; 1 = perfect alignment.
- **Normalized Mutual Information (NMI):** Measures shared information between predicted and true labels; 0 to 1.
- **Fowlkes-Mallows Index (FMI):** Geometric mean of precision and recall between cluster and label assignments; higher = better.

### 3. Generalizability / Stability

- Evaluate cluster consistency across variations of the dataset.
- Stable models produce similar clusters under small perturbations.

### 4. Dimensionality Reduction Evaluation

- **Explained Variance Ratio (PCA):** Amount of variance captured by components.
- **Reconstruction Error:** How accurately original data is reconstructed from reduced representation; lower = better.

- **Neighborhood Preservation:** How well high-dimensional relationships are maintained in lower dimensions (important for t-SNE, UMAP).

## 5. Visualization

- Scatter plots, dendograms, and projections (PCA, t-SNE, UMAP) are essential to interpret results.
  - Example: PCA on the Iris dataset shows PC1 and PC2 effectively separating species clusters.
- 

## 3. Practical Examples

- **K-means on Simulated Blobs:**
    - Distinct, dense clusters: Silhouette = 0.84, Davies-Bouldin = 0.22 → excellent clustering.
    - Somewhat dispersed clusters: Silhouette = 0.58, Davies-Bouldin = 0.6 → reasonable clustering.
  - **Dimensionality Reduction:**
    - First two principal components capture most variance in the Iris dataset, allowing visualization in 2D.
- 

## 4. Key Takeaways

- Unsupervised evaluation relies on **a combination of metrics, heuristics, domain expertise, and visualization.**
- **Internal metrics** assess cluster compactness and separation.
- **External metrics** compare clusters to true labels if available.
- **Dimensionality reduction** evaluation ensures important information is preserved.
- Stability and consistency are critical to trust the results of unsupervised models.

# Cross-Validation and Advanced Model Validation Techniques

## 1. What is Model Validation?

- Goal: Make sure your model **works well on unseen data** and isn't just memorizing the training data (avoiding overfitting).
  - Overfitting happens when a model performs well on the training data but poorly on new data.
  - **Data snooping / leakage:** Testing your model on the test set while tuning it is **cheating**, because the model indirectly "sees" the test data. This gives overly optimistic results.
- 

## 2. How to properly validate a model

- Split your data into **three parts**:

1. **Training set:** Used to train the model and adjust hyperparameters.
  2. **Validation set:** Used to evaluate different hyperparameter choices **during training**.
  3. **Test set: Held out** until the very end to check the final model's performance.
- 

### 3. Cross-Validation

- Standard train-validation split can be biased if you happen to pick a bad validation set.
  - **K-Fold Cross-Validation** helps:
    1. Split your training data into **K folds** (e.g., 5 or 10).
    2. For each fold:
      - Train on the other  $K-1$  folds
      - Validate on the current fold
    3. Repeat for all folds and **average the results**.
  - Benefits:
    - Every data point gets used for both training and validation.
    - Reduces overfitting to a specific validation set.
    - Gives a **more reliable estimate** of generalization performance.
- 

### 4. Stratified Cross-Validation

- Useful when your data is **imbalanced** (e.g., more examples of one class than another).
  - Ensures **class proportions are preserved** in each fold.
  - Prevents biased evaluation.
- 

### 5. Handling Skewed Targets (Regression)

- Some regression models assume a **normally distributed target**.
  - If the target is skewed, use **transformations** like:
    - **Log transform**
    - **Box-Cox transform**
  - These reduce skewness, helping the model fit better.
- 

### 6. Summary

- **Model validation prevents overfitting.**
- **Data snooping** = testing too early → avoid it.

- **Validation strategy:** training set + validation set + test set.
- **K-Fold cross-validation:** more robust, uses all data points for training and validation.
- **Stratified CV:** preserves class distributions for classification problems.
- **Target transformations:** handle skewed data for better regression results.

## Regularization in Regression and Classification

Regularization is used to **prevent overfitting** by adding a penalty to the size of the model's coefficients.

- constrains the model during training, discouraging it from overfitting to the training data.

It modifies the cost function (formula That tells how wrong its prediction is):

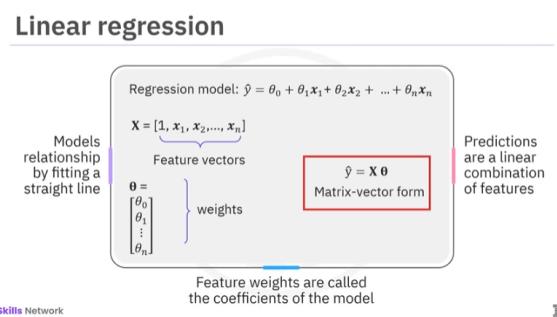
$$\text{Cost} = \text{MSE} + \lambda \times \text{Penalty}$$

where lambda: regularization hyperparameter

- Lambda is literally the knob that controls how much you punish large weights

penalty: Ridge, Lasso and other methods

Two main types:



## Ridge Regression (L2 penalty)

- Penalty = sum of squares of coefficients
- Shrinks coefficients **but never to zero**
- Useful when **all features are important**
- Good in noisy datasets

## Ridge and lasso regression

Ridge regression adds an L<sub>2</sub> penalty:  $\lambda \|\theta\|_2 = \lambda \sum_{i=1}^N \theta_i^2$

Lasso regression adds an L<sub>1</sub> penalty:  $\lambda \|\theta\|_1 = \lambda \sum_{i=1}^N |\theta_i|$

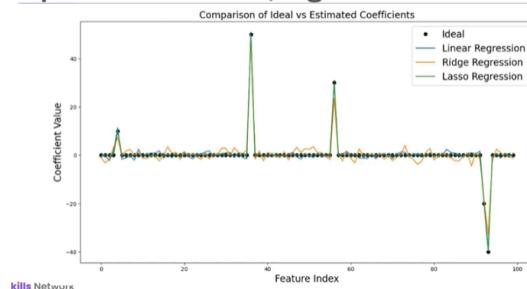
## Lasso Regression (L1 penalty)

- Penalty = sum of absolute values
- Shrinks some coefficients **exactly to zero**
-  Acts as **feature selection**
- Best when **only a few features matter (sparse)**
- Handles noise better than linear regression

## Linear Regression (no penalty)

- Performs well only when **data is clean (high SNR)**
- Fails badly when data is noisy
- Very sensitive to outliers

Sparse coefficients, high SNR

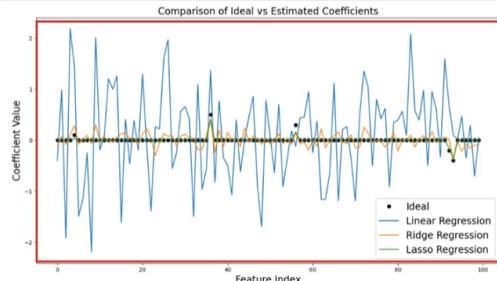


## Understanding SNR (Signal-to-Noise Ratio)

- **High SNR** → data is clean (less noise)
- **Low SNR** → data is noisy and messy

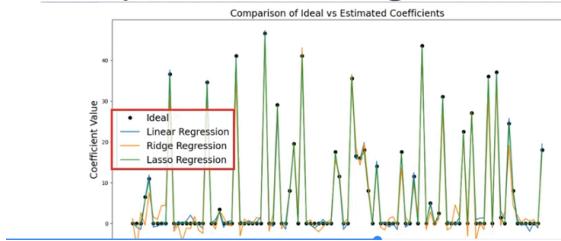
In low SNR:

### Sparse coefficients, low SNR



- Linear regression performs **very poorly**
- Ridge and Lasso are much better

### Non-sparse coefficients, high SNR



## 🎯 Main Takeaways

- Regularization helps control complexity and reduce overfitting.
- Lasso is best for sparse data and selecting features.
- Ridge is best when all features matter but need control.
- Linear regression only works well with clean, low-noise data.

## Data Leakage and Other Pitfalls

### 1. What is Data Leakage?

Data leakage happens when your model accidentally gets **information during training that it will not have in real life** after deployment.

Example:

You create a feature using the *average house price of the entire dataset*.

But in real life, when predicting a new house, you won't know this global average from the future data  
→ so the model performs unrealistically well.

This makes training/test accuracy look high, but in real-world use, the model fails.

### 2. Why is Leakage Dangerous?

- Your model **cheats** by learning from information it shouldn't know.
  - Test set also contains leaked info → test results look good.
  - When deployed, it performs badly because this leaked info won't exist.
- 

### 3. What Causes Data Leakage?

- Using **future information** to predict the present (e.g., using tomorrow's stock price).
  - Feature engineering using the **entire dataset** (e.g., global averages).
  - Poor train-test splitting (mixing test info into training).
  - Incorrect cross-validation where preprocessing is done before splitting.
- 

### 4. How to Avoid Leakage?

- Keep strict separation between **train**, **validation**, and **test** sets.
  - Avoid using features that depend on **future data**.
  - In pipelines, ensure preprocessing happens **inside** the CV loop for every fold.
  - For time-series data, use **TimeSeriesSplit** (not random splitting).
- 

### 5. Example Pipeline

A correct pipeline:

- Split data → define pipeline (scaler → PCA → KNN) → grid search → fit on train folds only → evaluate on held-out test set.

Key point:

**Pipeline is inside GridSearchCV** → ensures no leakage across folds.

#### Handling cross-validation data leakage

```
_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
stratify=y)
# Pipeline: StandardScaler, PCA, and KNN
pipeline = Pipeline([('scaler', StandardScaler()),
                    ('pca', PCA()),
                    ('knn', KNeighborsClassifier())
])
# Hyperparameter search grid for PCA components, KNN neighbors
param_grid = {'pca__n_components': [2, 3],
              'knn__n_neighbors': [3, 5, 7]
}
# Cross-validation method
cv = StratifiedKFold(n_splits=5, shuffle=True)
```

```

# Get best parameters
best_model = GridSearchCV(pipeline, param_grid, cv=cv, scoring='accuracy')

# Fit GridSearchCV to training data
best_model.fit(X_train, y_train)

# Evaluate the best model on the test set
test_score = best_model.score(X_test, y_test)

# Get the best parameters and score from grid search
print("Best Parameters:", best_model.best_params_)
print("Best Cross-Validation Accuracy:", best_model.best_score_)

# Evaluate the best model on the test set
print("Test Set Accuracy with Best Parameters:", best_model.score(X_test, y_test))
Best Parameters: {'knn__n_neighbors': 3, 'pca__n_components': 3}
Best Cross-Validation Accuracy: 0.962
Test Set Accuracy with Best Parameters: 0.933

```

## 6. Time-Series Data

For time-dependent data:

- You cannot randomly shuffle.
- Must split sequentially (train → older data, test → newer data).
- Use **TimeSeriesSplit** for cross-validation.

## 7. Feature Importance Pitfalls

Understanding feature importance is tricky. Problems include:

- **Redundant features share importance** → each looks less important.
- **Feature scale** affects models like linear regression.
- Importance ≠ causation → important feature does NOT mean it *causes* the outcome.
- Some models ignore **interactions** between features.

Example:

Two features individually weak, but their product is very powerful → linear regression fails, random forest succeeds.

## 8. Other Modeling Pitfalls

- Using raw data without cleaning/feature engineering.
- Wrong evaluation metric.
- Ignoring class imbalance.
- Blind trust in AutoML tools.
- Doing “what-if” predictions when the model has **no causal features** → results meaningless.

## Short Summary (Very Simple)

- **Data leakage** = model sees info during training that it won't see in real life.
- Leakage → model looks great in training/test but fails in production.
- Avoid leakage by:

- Proper train/validation/test split
- Avoid using future data
- Use pipelines inside cross-validation
- Use **TimeSeriesSplit** for time-based data
- Feature importance pitfalls:
  - Redundant features, scaling issues, correlation ≠ causation, ignored interactions.
- Modeling pitfalls:
  - Wrong metrics, ignoring imbalance, using raw data, trusting AutoML blindly.

## Model evaluation metrics and methods

| Method Name             | Description  | Code Syntax   |
|-------------------------|--|---|
| classification_report   | Generates a report with precision, recall, F1-score, and support for each class in classification problems. Useful for model evaluation. <b>Hyperparameters:</b> target_names: List of labels to include in the report. <b>Pros:</b> Provides a comprehensive evaluation of classification models. <b>Limitations:</b> May not provide enough insight for imbalanced datasets. | <pre> 1. from sklearn.metrics import classification_report 2. # y_true: True labels 3. # y_pred: Predicted labels 4. # target_names: List of target class names 5. report = classification_report(y_true, y_pred, target_names=["class1", "class2"]) Copied!Wrap Toggled!</pre> |
| confusion_matrix        | Computes a confusion matrix to evaluate the classification performance, showing counts of true positives, false positives, true negatives, and false negatives. <b>Hyperparameters:</b> labels: List of class labels to include. <b>Pros:</b> Essential for understanding classification errors. <b>Limitations:</b> Doesn't give insights into prediction probabilities.      | <pre> 1. from sklearn.metrics import confusion_matrix 2. # y_true: True labels 3. # y_pred: Predicted labels 4. conf_matrix = confusion_matrix(y_true, y_pred) Copied!Wrap Toggled!</pre>   |
| mean_squared_error      | Calculates the mean squared error (MSE), a common metric for regression models. Lower values indicate better performance. <b>Hyperparameters:</b> sample_weight: Weights to apply to each sample. <b>Pros:</b> Simple and widely used metric. <b>Limitations:</b> Sensitive to outliers, as large errors are squared.  | <pre> 1. from sklearn.metrics import mean_squared_error 2. # y_true: True values 3. # y_pred: Predicted values 4. # sample_weight: Optional, array of sample weights 5. mse = mean_squared_error(y_true, y_pred) Copied!Wrap Toggled!</pre>                                     |
| root_mean_squared_error | Calculates the root mean squared error (RMSE), which is the square root of the MSE. RMSE gives more interpretable results as it is in the same units as the target. <b>Hyperparameters:</b> sample_weight: Weights to apply to each sample. <b>Pros:</b> More  | <pre> 1. from sklearn.metrics import root_mean_squared_error 2. # y_true: True values 3. # y_pred: Predicted values 4. # sample_weight: Optional, array of sample weights 5. rmse =</pre>   |

| Method Name          | Description   | Code Syntax  |
|----------------------|---|--|
|                      | interpretable than MSE. <b>Limitations:</b> Like MSE, it can be sensitive to large errors and outliers.   | root_mean_squared_error(y_true,<br>y_pred)Copied!Wrap Toggled!   |
| mean_absolute_error  | Measures the average magnitude of errors in predictions, without considering their direction. Useful for understanding the average error size. <b>Hyperparameters:</b> sample_weight: Optional sample weights. <b>Pros:</b> Less sensitive to outliers compared to MSE. <b>Limitations:</b> Does not penalize large errors as much as MSE or RMSE.  | 1. from sklearn.metrics import<br>mean_absolute_error<br>2. # y_true: True values<br>3. # y_pred: Predicted values<br>4. mae =<br>mean_absolute_error(y_true,<br>y_pred)Copied!Wrap Toggled!                                       |
| r2_score             | Computes the coefficient of determination ( $R^2$ ), which represents the proportion of variance explained by the model. A higher value indicates a better fit. <b>Pros:</b> Provides a clear indication of model performance. <b>Limitations:</b> Doesn't always represent model quality, especially for non-linear models.                        | 1. from sklearn.metrics import<br>r2_score<br>2. # y_true: True values<br>3. # y_pred: Predicted values<br>4. r2 = r2_score(y_true,<br>y_pred)Copied!Wrap Toggled!   |
| silhouette_score     | Measures the quality of clustering by assessing the cohesion within clusters and separation between clusters. Higher scores indicate better clustering. <b>Hyperparameters:</b> metric: Distance metric to use. <b>Pros:</b> Useful for validating clustering performance. <b>Limitations:</b> Sensitive to outliers and choice of distance metric. | 1. from sklearn.metrics import<br>silhouette_score<br>2. # X: Data used in clustering<br>3. # labels: Cluster labels for each sample<br>4. score = silhouette_score(X,<br>labels,<br>metric='euclidean')Copied!Wrap Toggled!       |
| silhouette_samples   | Provides silhouette scores for each individual sample, indicating how well it fits its assigned cluster. <b>Hyperparameters:</b> metric: Distance metric to use. <b>Pros:</b> Offers granular insight into each sample's clustering quality. <b>Limitations:</b> Same as silhouette_score; sensitive to outliers and distance metric.               | 1. from sklearn.metrics import<br>silhouette_samples<br>2. # X: Data used in clustering<br>3. # labels: Cluster labels for each sample<br>4. samples =<br>silhouette_samples(X, labels,<br>metric='euclidean')Copied!Wrap Toggled! |
| davies_bouldin_score | Measures the average similarity ratio of each cluster with the most similar cluster. Lower values indicate better clustering. <b>Pros:</b> Provides a simple, effective clustering evaluation. <b>Limitations:</b> May not work well with highly imbalanced clusters.   | 1. from sklearn.metrics import<br>davies_bouldin_score<br>2. # X: Data used in clustering<br>3. # labels: Cluster labels for each sample<br>4. db_score =<br>davies_bouldin_score(X,<br>labels)Copied!Wrap Toggled!                |
| Voronoi              | Computes the Voronoi diagram, which partitions space based on the nearest neighbor. <b>Pros:</b> Useful for spatial analysis and clustering. <b>Limitations:</b> Limited to use cases that involve spatial partitioning of data.  | 1. from scipy.spatial import<br>Voronoi<br>2. # points: Coordinates for<br>Voronoi diagram<br>3. vor =   |

| Method Name              | Description   | Code Syntax   |
|--------------------------|---|---|
|                          |   | Voronoi(points)Copied!Wrap<br>Toggled!  |
| voronoi_plot_2d          | Plots the Voronoi diagram in 2D for visualizing clustering.<br><b>results.Hyperparameters:</b> show_vertices: Whether to display the vertices. <b>Pros:</b> Great for visualizing spatial clustering. <b>Limitations:</b> Limited to 2D spaces and large datasets may cause performance issues.                     | 1. from scipy.spatial import voronoi_plot_2d<br>2. # vor: Voronoi diagram object<br>3. voronoi_plot_2d(vor, show_vertices=True)Copied!Wrap<br>Toggled!  |
| matplotlib.patches.Patch | Creates custom shapes such as rectangles, circles, or ellipses for adding to plots. <b>Hyperparameters:</b> color: Fills color of the shape. <b>Pros:</b> Versatile for visual customization. <b>Limitations:</b> May not support all shapes or complex customizations.   | 1. import matplotlib.patches as patches<br>2. # Create a rectangle with specified width, height, and position<br>3. rectangle = patches.Rectangle((0, 0), 1, 1, color='blue')Copied!Wrap<br>Toggled!  |
| explained_variance_score | Measures the proportion of variance explained by the model's predictions. A higher score indicates better performance. <b>Pros:</b> Helps in assessing the fit of regression models. <b>Limitations:</b> Not suitable for classification tasks.   | 1. from sklearn.metrics import explained_variance_score<br>2. # y_true: True values<br>3. # y_pred: Predicted values<br>4. ev_score = explained_variance_score(y_true, y_pred)Copied!Wrap<br>Toggled! |
| Ridge regression         | Performs ridge regression (L2 regularization) to avoid overfitting by penalizing large coefficients. <b>Hyperparameters:</b> alpha: Regularization strength. <b>Pros:</b> Helps reduce overfitting in regression models. <b>Limitations:</b> May not work well with sparse data.                                    | 1. from sklearn.linear_model import Ridge<br>2. # alpha: Regularization strength (larger values indicate stronger regularization)<br>3. ridge = Ridge(alpha=1.0)Copied!Wrap<br>Toggled!               |
| Lasso regression         | Performs lasso regression (L1 regularization), which encourages sparsity by penalizing the absolute value of coefficients. <b>Hyperparameters:</b> alpha: Regularization strength. <b>Pros:</b> Encourages sparse solutions, useful for feature selection. <b>Limitations:</b> May struggle with multicollinearity. | 1. from sklearn.linear_model import Lasso<br>2. # alpha: Regularization strength (larger values indicate stronger regularization)<br>3. lasso = Lasso(alpha=0.1)Copied!Wrap<br>Toggled!               |
| Pipeline                 | Chains multiple steps of preprocessing and modeling into a single object, ensuring efficient workflow. <b>Pros:</b> Simplifies code, ensures reproducibility. <b>Limitations:</b> May not work well with complex pipelines requiring dynamic configurations.  | 1. from sklearn.pipeline import Pipeline<br>2. # steps: List of tuples with name and estimator/transformer<br>3. pipeline = Pipeline(steps=[('scaler', StandardScaler()), ('model',                   |

| Method Name  | Description  | Code Syntax  |
|--------------|--|--|
| GridSearchCV | <p>Performs exhaustive search over a specified parameter grid to find the best model configuration. <b>Hyperparameters:</b> param_grid: Dictionary of parameter grids. <b>Pros:</b> Ensures optimal model parameters. <b>Limitations:</b> Computationally expensive for large grids.</p> | <pre>Ridge(alpha=1.0)))])Copied!Wrap Toggled!</pre> <pre>1. from sklearn.model_selection import GridSearchCV 2. # estimator: Model to be tuned 3. # param_grid: Dictionary with parameters to search over 4. grid_search = GridSearchCV(estimator=Ridge(), param_grid={'alpha': [0.1, 1.0, 10.0]})Copied!Wrap Toggled!</pre> |

## Visualization strategies for k-means evaluation

| Process Name             | Brief Description   | Code Snippet   |
|--------------------------|---|--|
| Multiple runs of k-means | <p>Executes KMeans clustering multiple times with different random initializations to assess variability in cluster assignments. <b>Advantage:</b> Helps visualize consistency. <b>Limitation:</b> Computationally costly for large datasets.</p> | <pre>1. # Number of runs for KMeans with different random states 2. n_runs = 4 3. inertia_values = [] 4. 5. plt.figure(figsize=(12, 12)) 6. 7. # Run K-Means multiple times with different random states 8. for i in range(n_runs): 9.     kmeans = KMeans(n_clusters=4, random_state=None) # Use the default `n_init` 10.    kmeans.fit(X) 11.    inertia_values.append(kmeans.inertia_) 12. 13. # Plot the clustering result 14. plt.subplot(2, 2, i + 1) 15. plt.scatter(X[:, 0], X[:, 1], c=kmeans.labels_, cmap='tab10', alpha=0.6, edgecolor='k') 16. plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], c='red', s=200, marker='x', label='Centroids') 17. plt.title(f'K-Means Clustering Run {i + 1}') 18. plt.xlabel('Feature 1') 19. plt.ylabel('Feature 2') 20. plt.legend() 21. 22. plt.tight_layout() 23. plt.show() 24. 25. # Print inertia values 26. for i, inertia in enumerate(inertia_values, start=1): 27.     print(f'Run {i}: Inertia={inertia:.2f}')</pre> |
| Elbow method             | <p>Evaluates the optimal number of clusters by plotting inertia (within-cluster sum of squares) for different k values. <b>Advantage:</b> Easy to</p>   | <pre>1. # Range of k values to test 2. k_values = range(2, 11) 3. 4. # Store performance metrics</pre>   |

| Process Name                | Brief Description   | Code Snippet  |
|-----------------------------|---|---|
|                             | interpret. <b>Limitation:</b> Subjective elbow point.   | <pre> 5. inertia_values = [] 6. for k in K_values: 7.     kmeans = KMeans(n_clusters=k, random_state=42) 8.     y_kmeans = kmeans.fit_predict(X) 9. 10.    # Calculate and store metrics 11.    inertia_values.append(kmeans.inertia_) 12. 13.    # Plot the inertia values (Elbow Method) 14.    plt.figure(figsize=(18, 6)) 15.    plt.subplot(1, 3, 1) 16.    plt.plot(k_values, inertia_values, marker='o') 17.    plt.title('Elbow Method: Inertia vs. k') 18.    plt.xlabel('Number of Clusters (k)') 19.    plt.ylabel('Inertia')Copied!Wrap Toggled! </pre>   |
| <b>Silhouette method</b>    | <p>Determines the optimal number of clusters by evaluating Silhouette Scores for different <b>k</b> values. <b>Advantage:</b> Considers both cohesion and separation. <b>Limitation:</b> High computation for large datasets.</p> | <pre> 1. # Range of k values to test 2. k_values = range(2, 11) 3. 4. # Store performance metrics 5. silhouette_scores = [] 6. for k in k_values: 7.     kmeans = KMeans(n_clusters=k, random_state=42) 8.     y_kmeans = kmeans.fit_predict(X) 9.     silhouette_scores.append(silhouette_score(X, y_kmeans)) 10. 11. # Plot the Silhouette Scores 12. plt.figure(figsize=(18, 6)) 13. plt.subplot(1, 3, 2) 14. plt.plot(k_values, silhouette_scores, marker='o') 15. plt.title('Silhouette Score vs. k') 16. plt.xlabel('Number of Clusters (k)') 17. plt.ylabel('Silhouette Score')Copied!Wrap Toggled! </pre>                           |
| <b>Davies-Bouldin Index</b> | <p>Evaluates clustering performance by calculating DBI for different <b>k</b> values. <b>Advantage:</b> Quantifies compactness and separation. <b>Limitation:</b> Sensitive to cluster shapes and density.</p>                    | <pre> 1. # Range of k values to test 2. k_values = range(2, 11) 3. 4. # Store performance metrics 5. davies_bouldin_indices = [] 6. for k in k_values: 7.     kmeans = KMeans(n_clusters=k, random_state=42) 8.     y_kmeans = kmeans.fit_predict(X) 9. davies_bouldin_indices.append(davies_bouldin_score(X, y_kmeans)) 10. 11. # Plot the Davies-Bouldin Index 12. plt.figure(figsize=(18, 6)) 13. plt.subplot(1, 3, 3) 14. plt.plot(k_values, davies_bouldin_indices, marker='o') 15. plt.title('Davies-Bouldin Index vs. k') 16. plt.xlabel('Number of Clusters (k)') 17. plt.ylabel('Davies-Bouldin Index')Copied!Wrap Toggled! </pre> |